

Graphs

Subhash Suri

April 25, 2017

1 Basic Definitions

- Graphs are useful models for reasoning about relations among objects and combinatorial problems. Many real-life problems can be solved by converting them to graphs. Proper application of graph theory ideas can drastically reduce the solution time for some important problems.
- A graph has a set vertices V , often labeled v_1, v_2, \dots , and a set of edges E , labeled e_1, e_2, \dots .
- Each edge (u, v) “joins” two nodes u and v .
- We write $G = (V, E)$ for the graph with vertex set V and edge set E .
- In applications, where pair (u, v) is distinct from pair (v, u) , the graph is *directed*. Otherwise, the graph is undirected. We can convert an undirected graph to a directed one by duplicating edges, and orienting them both ways.
- When (u, v) is an edge, we say v is *adjacent to (or, neighbor of) u* . A loop is an edge with both endpoints being the same.
- In undirected graphs, the *degree* of a node equals its number of neighbors. In directed graphs, we have The *out-degree* and the *in-degree*.
- In some applications, the edges can be associated with weights or costs.

2 Examples of Graphs

- Transportation Networks. The map of routes served by an airline carrier forms a graph, whose nodes are the airports, and we have an edge (u, v) whenever airline has a non-stop flight from u to v . Typically, airline edges are undirected—flight (u, v) also means a flight (v, u) .

Other transportation networks: rail networks, road networks.

- Communication Networks. Internet is essentially a collection of computers connected by communication links. Nodes are computers, and edges are physical links.

Wireless networks: devices, and wireless connections.

- Information networks. WWW has web pages as nodes, and hyperlinks as edges.
- Social networks.
- Dependency graphs: nodes = courses, and edges = prereqs;

3 Representations of Graphs

- **Adjacency Matrix:** a 2-dim array $V \times V$. For each edge (u, v) , set $A[u, v] = 1$, or equal to cost, etc. Use infinity or 0 for non-edges.
- Pros: easy to check if (u, v) an edge in G .
- Cons: Takes V^2 space even if graph has very few edges; e.g. street map, which typically has $O(V)$ edges. Infeasible space when V is millions of nodes.
- **Adjacency List:** An array of (header cells for) adjacency lists. The i th cell points to a linked list of all vertices adjacent to vertex v_i .
- Example:

1 :	2	4	3
2 :	4	5	
3 :	6		
4 :	6	7	3
5 :	4	7	
6 :			
7 :	6		

- Space is $O(E)$; each directed edge stored just once. Thus, if G is undirected (u, v) appears in lists of both u and v .
- Pros. Linear space. Easy to list out all vertices adjacent to u .
- Cons: Checking if (u, v) is an edge can take $O(n)$ time.

4 Paths and Connectivity

- One of the fundamental operations in graphs is that of traversing a sequence of nodes (and edges). Such a traversal could correspond to user browsing web pages by following hyper links, rumor passing by word of mouth, or travel route of an airline passenger, email passing through a chain of routers, etc.
- A path is sequence of vertices w_1, w_2, \dots, w_k such that each pair (w_i, w_{i+1}) is an edge of G . The *length* of a path is the number of edges in it, or total weight if each edge has a weight associated with it.
- A simple path has no repeated vertex, except first and last can be the same; in that case, the path is a cycle.
- An undirected graph is *connected* if there is a path between any two vertices. A directed graph with this property is *strongly connected*. A weakly connected graph—underlying graph connected but the directed graph may not have directed path between all pairs.
- **Trees:** an undirected graph is a tree if it is connected and does not contain a cycle.
- Trees are one of the simplest type of graphs. Any tree on n nodes has $n - 1$ edges, and therefore the deletion of a single node or edge disconnects it.
- Often, it is useful to *root* the tree at a particular node r , and then *orient* all edges away from r . EXAMPLE.
- In a rooted tree, each node (except root) has a parent, and if u is the parent of w , then w is called a *child* of u .
- More generally, w is called a *descendant* of u , and u an *ancestor* of w , if u lies on the path from w to the root.

5 Graph Connectivity and Graph Traversals

- We start with one of the most basic questions regarding graphs. Given a graph $G = (V, E)$, and two nodes s and t , is there a path joining s and t ?
- This is called the *st-connectivity* problem. (This is also the classical Maze problem.) In small graphs, one can decide this by visual inspection, but quickly becomes challenging in large graphs.
- More generally, given a start node s , what are all the nodes reachable from s ? This set is called the *connected component* of G containing s .

- There are two simply algorithms for st -connectivity

5.1 Breadth First Search

- The simplest algorithm for st -connectivity is the following. We start at s , and explore outward in all possible directions.
- We just have to make sure we don't get stuck in a loop, so we use *markers* to keep track of nodes we have already visited.
- Each node will get a *layer number* (also called level). Initially, we have only s , which is layer 0. The next iteration adds previously unreached nodes that have an edge to an already reached node. More specifically,
 - We initialize Layer $L_0 = \{s\}$; i.e. layer 0 containing just s . Layer L_1 consists of all neighbors of s .
 - Assuming we have layers L_0, L_1, \dots, L_i , define

$$L_{i+1} = \text{nodes not yet encountered who have an edge to some node in layer } L_i$$

- Example with 3 connected components.
- The layer by layer exploration of G produces a tree-like structure, which is called the BFS tree of G .
- For each $j \geq 1$, layer L_j consists of all nodes at distance exactly j from s . There is a path from s to t if and only if s appears in some layer of BSF from s .
- Let T be a BFS tree, and let x and y be nodes in T belonging to different layers L_i and L_j . If (x, y) is an edge of G , then $|i - j| \leq 1$.
- Proof. For contradiction, assume $i < j - 1$. When x is scanned in layer i , the edge (x, y) will add y to layer L_{i+1} , ensuring $j \leq i + 1$.
- BFS can be constructed in $O(m + n)$ time, using Adj List representation of G .
- The set of nodes discovered by the BFS is precisely those reachable from s . We refer to this set R as the *connected component* of G containing s .
- Once we have R , we can simply check if t belongs to R , and if so we have st connectivity.
- BFS however is only one way to discover R . Another, and a very different, method is depth first search.

6 Depth First Search

- Depth-First-Search (DFS) uses a method similar to the exploration of mazes:
- Starting at s , we take the first edge out of s , and continue recursively until we reach a *dead end*—a node for which all neighbors have already been explored.
- We then backtrack until we get to a node with at least one explored neighbor.
- This is called DFS search, because it explores G by going as deeply as possible and only retreating when necessary.

```
DFS(u)
    Mark u as Explored and add u to R
    For each edge (u, v) incident to u
        if v is not marked Explored, then
            recursively call DFS(v)
        endif
    endfor.
```

- We can also implement DFS non-recursively.

Stack Implementation of DFS:

```
DFS(s)
    Init S to be a stack with one item s
    While S not empty
        Take a node u from S
        If Explored[u] = False then
            Set Explored[u] = True
            For each edge (u,v) incident to u
                Add v to stack S
            endfor
        endif
    endwhile
```

- Example from Kleinberg-Tardos.
- DFS also runs in $O(m + n)$ time, where $n = |V|$ and $m = |E|$.
- Although the DFS tree looks very different from the BFS tree of G , we can make strong claims about how non-tree edges connect the nodes of DFS.
- **Fact 1.** For a recursive call $DFS(u)$, all nodes that are marked *explored* between the invocation and the end of the recursive call are descendants of u in T .
- **Fact 2.** Let T be a DFS tree, let x, y be two nodes in T that have an edge between them in G , but (x, y) is not an edge of T . Then, one of x or y is an ancestor of the other.
- Suppose not, and assume that x is reached first in DFS. When the edge (x, y) is examined during the execution of $DFS(x)$, it is not added to T because y is marked Explored. Since y was not marked Explored when $DFS(x)$ was first invoked, it must have been discovered during the recursive call. Thus, by Fact 1, y must be a descendant of x .
- **Connected Components Fact.** For any two nodes s and t in G , their connected components are either identical or disjoint.

7 Applications of BFS and DFS

- **Testing Bipartiteness.** A graph G is bipartite if its vertex set V can be partitioned into sets X and Y in such a way that every edge of G has one end in X and the other in Y .
- Often we use colors red and blue (or 0 and 1) to represent X and Y .
- A triangle is not bipartite: any partition will contain two nodes on the same side with an edge between them. The same argument also holds if G is an odd-length cycle.
- Turns out however that odd cycles are the only obstacle for G to be bipartite: that is, G is bipartite if and only if it does not contain an odd cycle.
- In fact, one can use BFS to decide whether G is bipartite, and in the end either discover the sets X and Y , or detect an odd-cycle, thereby showing that G is not bipartite.
- We can easily assume that G is connected. Otherwise, we can apply the algorithm to each connected component separately.

- The algorithm begins by picking any arbitrary vertex s , and color it 0.
- Now all neighbors of s must be colored 1, and these are precisely the nodes of layer 1.
- We alternate between colors: the nodes at layer i are colored 0 if i is even, and colored 1 if i is odd.
- At the end of the algorithm, we simply go back and check if the endpoints of each edge of G are colored differently. If not, that edge (x, y) together with the path in the BFS from x to y is an odd cycle.
- Therefore, bipartiteness of a graph G can be decided in $O(m + n)$ time.

8 Bi-Connectivity

- An undirected graph G is *bi-connected* if the deletion of a single node keeps it connected. That is, one must delete at least two nodes (and their incident edges) to disconnect G .
- Another classical application of DFS is a linear-time algorithm (due to Hopcroft and Tarjan) to find bi-connected components of G .
- **Articulation point** is a node v whose removal disconnects G . Thus, G is bi-connected if and only if there is no articulation point.
- The main idea is to run a DFS while maintaining the following information for each vertex v of the DFS tree T :
 1. the depth of v (once it gets visited), and
 2. the lowest depth among the neighbors of all descendants of v , called the *lowpoint*
- More specifically, let $d(v)$ be the depth (DFS number) of node v . Define

$$low(v) = \min\{d(v), \{d(w) : (u, w) \text{ is a back edge for some descendant } u \text{ of } v\}\}$$

- The $low()$ values of all the nodes can be computed in linear time, by performing a *post-order* traversal of T .
- Example.
- Once we have these computed, detecting articulation points is easy: the root is an articulation point, if it has more than one child; any non-root node v is an articulation point if it has a child w with $low(w) \geq d(v)$.

- For proof, notice that if v is an articulation point then none of the nodes explored during the recursive call at v have an edge that goes to the other component, and thus the $low()$ value for all these points is $\geq d(v)$.

9 Topological Sort

- Suppose you have a set of tasks, which are subject to a set of precedence constraints: some jobs cannot be done before others. How shall you schedule the jobs without violating any prec constraint?
- Model as a *directed* graph where jobs are nodes and precedence relations are edges.
- Clearly, if there is a cycle in the graph, no feasible schedule.
- When there is no cycle, *topological sorting* is an ordering of vertices such if there is a path from v_i to v_j , then v_i appears *before* v_j in the schedule.

Algorithm:

```
Find a vertex v with zero in-degree (must exist!)
Print v, delete v, and its outgoing edges;
Repeat.
```

Improved Topological Sort

```
Compute all vertices' indegs
Enqueue all those with zero indeg
Pick a vertex w from the queue;
    put w next in schedule
    for each vertex v adj to w
        decrement v's indeg
        add v to queue if its indeg = 0
```

- This code only looks at each edge once, so $O(E)$ time.
- Example.
- One can use DFS to also perform topological sorting. How?

10 Strong Bi-Connectivity

- DFS and BFS algorithms work on directed graphs, without any significant change: while visiting a vertex v , we just scan v 's out neighbors.
- In directed graphs, however, we need a stronger definition of a connected components. We put two vertices u and v in the same component only if we have a directed path from u to v **and** a path from v to u .
- Example.
- We can also find strong connected components of G also in $O(|V| + |E|)$ time, by using DFS, but in a more careful way.
- Historically, the first linear time algorithm dates back to 70s by Hopcroft and Tarjan.
- A simpler algorithm is by Koraraju-Sharir. It performs two DFS once on G , and once on G^R , which is G with all edges reversed.
- Intuition. Perform DFS on G , and list the vertices in the **post-order**.
- Figure 1 shows a directed graph, and its DFS.




Figure 1: A directed graph and its DFS.

- The post-order numbering of nodes is: $G, H, J, I, B, F, C, A, D, E$.
- We now perform a DFS on G^R , always starting new DFS at the highest numbered vertex. So, in the example, first DFS starts at node G , numbered 10. This leads nowhere, so G is a singleton node component.
- See Figure 2.
- Next DFS starts at H , and this call adds I and J to the component of H .

- Next starts at B , and adds $\{A, C, F\}$ before finishing.
- DFS at D ends with singleton, as does for E .




Figure 2: G^R , with post-order numbering from the first DFS.

- Proof of Correctness. Key idea is that if u, v are in the same SCC, then there are paths from u to v , and from v to u , in both G and G^R .
- Thus, if two nodes are not in the same DFS tree, then they cannot be in one SCC.
- We show that if x is the root of the DFS tree in G^R containing v , then there is a path from x to v , and from v to x . Applying the same logic to w gives a pair of paths between x and w , and thus shows that x, v, w are in the same SCC.
- Since v is a descendant of x in G^R DFS, there is path from x to v in G^R , and thus a path from v to x in G .
- Since x is the root, it has the higher post-order than v . Therefore, during the DFS in G , the recursive call at v finished before the recursive call at x finished. Since a path from v to x exists, it must be that v is a descendant of x in the DFS of G —otherwise, v would finish *after* x . Therefore, there is a path from x to v , and the proof is complete.